


Reference:

[FAQs](#)

Category:

Compute



Amazon EC2

What?

- Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides resizable computing capacity—literally, servers in Amazon's data centers—that you use to build and host your software systems.
- An instance is a virtual server in the AWS Cloud.

Why?

- Amazon EC2 reduces the time required to obtain and boot new server instances to minutes, allowing you to quickly scale capacity, both up and down, as your computing requirements change.
- Amazon EC2 changes the economics of computing by allowing you to pay only for capacity that you actually use.

When?

- Run cloud-native and enterprise applications, Scale for HPC applications, Develop for Apple platforms.
- You want complete control of your computing resources and run it on Amazon's proven computing environment.
- You want to import your virtual machine images to Amazon EC2.

Where?

- Amazon EC2 is a regional service. An EC2 Instance runs in an Availability Zone.
- By launching instances in separate Availability Zones, you can protect your applications from failure of a single location.

Who?

- Customer must take care of OS patches, high availability and scaling. AWS provides tools and services for Patch Management, Auto Scaling, Monitoring, Backup and Vulnerability Scanning.

How?

- Select an AMI (Amazon Machine Image) >> Select Instance Type >> Select Additional Settings (Disk Size / Security Group Setting / Network / Start-up Script etc.).

How much?

- On-Demand Instance – Pay per second / hour.
- Reserved Instances / Savings Plan – Commitment for 1-year / 3-year term.
- Spot Instances – Supply and Demand based pricing. May be terminated by AWS after giving a 2 min notice.

Created by:

[Ashish Prajapati](#)

