Reference:

FAQs

Category:

Database

Amazon Aurora

What?

- Amazon Aurora is a MySQL and PostgreSQL-compatible relational database built for the cloud that combines the performance and availability of traditional enterprise databases with the simplicity and cost-effectiveness of open source databases.
- It features a distributed, fault-tolerant, and self-healing storage system that is decoupled from compute resources.

Why?

- Aurora automates time-consuming administration tasks like hardware provisioning, database setup, patching, and backups while providing the security, availability, and reliability of commercial databases at 1/10th the cost.

When?

- Amazon Aurora is a great option for any enterprise application that can use a relational database.
- You need high performance and availability with up to 15 low-latency read replicas, point-in-time recovery, continuous backup to Amazon S3, and replication across three AZs.

Where?

- Amazon Aurora is a regional service, it automatically maintains six copies of your data across three AZs.
- Cross-region Aurora replicas can be setup using either physical or logical replication. Physical replication uses Amazon Aurora Global Database, logical replication uses binlog for MySQL and PostgreSQL replication slots for PostgreSQL

Who?

- Amazon Aurora is fully managed by RDS and it automatically and continuously monitors and backs up your database to Amazon S3, enabling granular point-in-time recovery.
- Customer can scale the compute resources allocated to your DB Instance by changing your DB Instance class.

Complete book:

Click Here

How?

- You choose Aurora as the DB engine option when setting up new database servers through Amazon RDS.
- After launching an Aurora instance, you can connect to it using any database client that supports MySQL or PostgreSQL.

Created by:

Ashish Prajapati

How much?

- For provisioned Aurora, you can choose On-Demand Instances and pay for your database by the hour with no long-term commitments or upfront fees, or choose Reserved Instances for additional savings.
- Aurora storage is billed in per GB-month increments, while I/Os consumed are billed in per million request increments.

Reference:

FAQs

Category:

Database

Amazon Aurora Serverless

What?

- Amazon Aurora Serverless is an on-demand, autoscaling configuration for Amazon Aurora. It automatically starts up, shuts down, and scales capacity up or down based on your application's needs.
- Aurora Serverless v2 is available for Aurora MySQL-Compatible and PostgreSQL-Compatible editions.

Why?

- Manually managing database capacity can take up valuable time and can lead to inefficient use of database resources. With Aurora Serverless, you create a database, specify the desired database capacity range, and connect your applications.
- It supports the full breadth of Aurora features, including global database, Multi-AZ deployments, and read replicas.

When?

- You want to run your database on AWS provisioning and managing database capacity.
- You have intermittent, infrequent, or unpredictable bursts of requests and want your database to automatically scale capacity to meet the needs of the application's peak load and scale back down when the surge of activity is over.

Where?

- Aurora Serverless is a Regional service. A Multi-AZ Aurora DB cluster has compute capacity available more than one AZ.
- Amazon Aurora Global Database allows a single Amazon Aurora database to span multiple AWS Regions.
- The storage for each Aurora DB cluster consists of six copies of all your data, spread across three AZs.

Who?

- There is no database capacity for you to manage. Amazon Aurora Serverless automatically starts up, scales compute capacity to match your application's usage, and shuts down when it's not in use.
- You can upgrade or switch existing clusters to use Aurora Serverless v2.

Complete book:

Click Here

How?

- Create a database, specify the desired database capacity range (minimum and maximum amount of resources needed), and connect your application. Aurora automatically adjusts the capacity within the range based on your application’s needs.

Created by:

Ashish Prajapati

How much?

- Aurora Serverless measures database capacity in Aurora Capacity Units (ACUs) billed per second when the database is active.
- 1 ACU has approximately 2 GiB of memory with corresponding CPU and networking. Amazon Aurora database storage consumption is billed in per GB-month increments, and I/Os consumed are billed in per million request increments.